

Deployed Resources — Meridian Bank Retention & Cross-Sell Radar

Workspace: `fevm-vijay` ([https://fevm-vijay.cloud.databricks.com, o=7474659668600833](https://fevm-vijay.cloud.databricks.com,o=7474659668600833))

Catalog / schema: `vijay_catalog.banker_retention_cross_sell`

Deployed: 2026-08-27 (target `dev`)

Is there a Genie Agent?

Yes — but it is **not** a separately-hosted agent or model-serving endpoint. The agent lives *inside the Databricks App*: `app/server/agent/retentionops.ts` runs an in-process

LLM agent (via `@openai/agents`) that calls a foundation model through the **AI Gateway**,

and its single data tool is `ask_genie` (`app/server/agent/tools/genie.ts`), which drives

the Genie space through the `/api/2.0/genie/spaces/{id}` conversation API.

Topology: **App** → **agent loop** → `ask_genie` **tool** → **Genie space**.

There is **no** Knowledge Assistant, **no** Multi-Agent Supervisor, and **no** registered/served ML model (the `KA_ENDPOINT_NAME` / `MAS_ENDPOINT_NAME` env vars are empty). The only AI resource

object in the workspace is the Genie space; the agent itself ships as app code.

Key resources

Application & AI

1. **Databricks App — `dbgen-meridian-radar`** (`manage`)

The RM-facing web app: customer 360, next-best-action, and draft retention/cross-sell offers with write-back; hosts the Genie-backed agent.

2. **Genie space — Meridian Retention & Cross-Sell**

The natural-language analytics space (over the gold tables + metric view) that the app's `ask_genie` tool queries.

Analytics

1. **AI/BI Dashboard**

The book-of-business dashboard: at-risk balances, attrition trends, and cross-sell opportunity by segment/branch.

2. **Metric view — mv_book_health**

Governed metric layer (MEASUREs like `at_risk_balance`, `attrition_risk_rate`, `cross_sell_opportunity`) used by both Genie and the dashboard.

Data pipeline & orchestration

1. **Lakeflow pipeline — Meridian Retention Radar Pipeline**

Serverless + Photon SDP pipeline that builds silver → gold from the raw parquet.

2. **Setup job — Meridian Retention Radar Setup**

The one-shot orchestrator: generate data → run pipeline → deploy metric view + Genie → grant app SP → export IDs.

Unity Catalog data

1. **Schema — vijay_catalog.banker_retention_cross_sell**

The demo's home schema holding all tables, views, and the volume.

2. **Volume — raw_data**

Raw parquet landing zone (customers, accounts, transactions, balance_weekly).

3. **gold_customer_360**

Per-customer risk scores, runoff %, payroll signals, and next-best-action — the app's core read model.

4. **gold_rm_radar & gold_weekly_book_summary**

RM-level radar and weekly book rollups. (Plus 4 `silver_*` views and 4 `_`-prefixed intermediate materialized views.)

Storage & observability

(No direct deep-link — identifiers given.)

1. **Lakebase (Postgres) database — dbgen_meridian_radar**

Project `dbdemos-asset-generator`, branch `production`, host `ep-fragrant-haze-d214556c.database.us-east-1.cloud.databricks.com`; the app's transactional store for offers / write-back. Find it under **Compute → Database instances**.

2. **MLflow experiment — /Shared/solution_builder/dbgen-meridian-radar-agent-traces**

Auto-created at app boot to capture the agent's traces; browse under **Experiments**.

Resolved IDs (from the setup job's `export_resources`)

Resource	ID
Catalog	vijay_catalog
Schema	banker_retention_cross_sell
App	dbgen-meridian-radar
Genie space	01f1a25bc3481bd78367ba5c93c7a886
Dashboard	01f1a254cd2f1144b6d2e1f105431e2d
Pipeline	bca0a467-3418-4a95-894b-93496673db82
Setup job	947855350108729
SQL warehouse	79fae613dbebe6d9
Agent MLflow experiment path	/Shared/solution_builder/dbgen-meridian-radar- agent-traces